

MC 302 – DBMS: Hashing

Goonjan Jain

Department of Applied Mathematics

Delhi Technological University

Outline

- hashing
- extendible hashing
- linear hashing
- Hashing vs B-trees

Hashing

- Primary file organization
- Very fast access to records on search condition
- Search condition must be on a single field – hash field or hash key
- Hash function/ randomizing function – applied on hash field, yields address of disk block

Static Hashing


- Fixed number of buckets
- Drawback for dynamic files
- Types:
 - Internal Hashing
 - External Hashing

Internal Hashing

- For internal files
- implement a hash table using array of records
- M slots addressed as 0 to M-1
- Choose a hash function, transform hash field into an integer from 0 to M-1
- Most common hash function $h(k) = k \bmod M$
- Problems:
 - No guarantee that different values will hash to different addresses

Collision

- Hash field value hashed to an address already occupied
- Collision resolution: find another position
- Collision Resolution Techniques:
 - Open addressing:
 - Starting from the hashed address, check subsequent positions till an unused position is found
 - Chaining:
 - Place new value in an unused overflow location
 - Set a pointer from the address to the overflow location



- Multiple Hashing:
 - Apply a second hash function

External Hashing


- Target address space is made of **buckets**,
 - each bucket can hold multiple values
- Bucket is
 - 1 disk block, or
 - Cluster of contiguous blocks
- Hash function maps key to a relative bucket number
- Collision problem is less severe
- If a bucket is full, a variation of chaining can be used-
 - Bucket points to record pointers – block address and record position

Problem with (static) hashing

- Overflow
- Underflow
- Sol: Dynamic Hashing

Dynamic Hashing


- idea: shrink / expand hash table on demand..
- ..dynamic hashing
- Details: how to grow gracefully, on overflow?
- Many solutions – One of them: ‘**extendible hashing**’



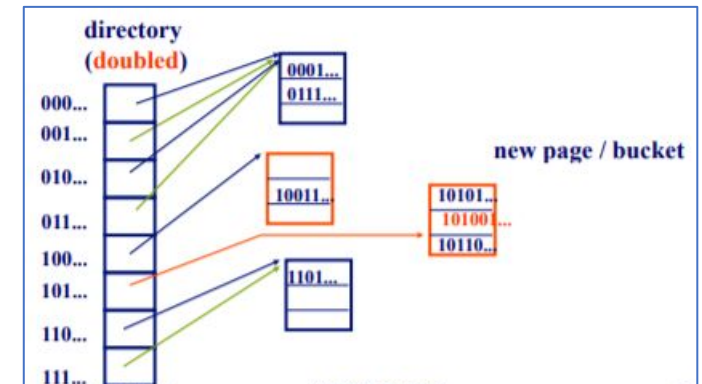


Extendible Hashing

- keep a directory, with pointers to hash-buckets
- Uses-
 - An array of 2^d buckets. d is the global depth
 - Directory
- First d bits of hash value is used as index for directory entry
- Entry in director determines the bucket address
- Each bucket stores
 - Local depth d' – number of bits on which bucket contents are based
- Q: how to divide contents of bucket in two?
- A: hash each key into a very long bit string; keep only as many bits as needed


Extendible Hashing



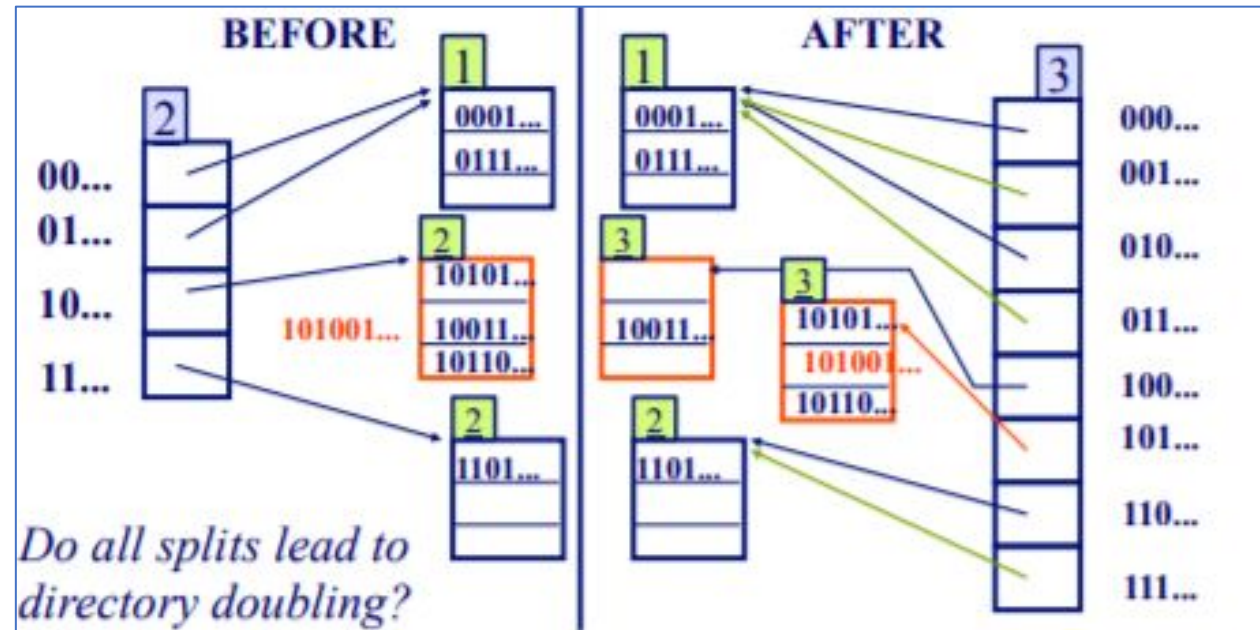
Extendible Hashing



Extendible Hashing



Extendible Hashing



Extendible Hashing – Directory Doubling

- bucket overflows, 2 cases –
 - If local depth, d' = global depth, d
 - Double the directory
 - If $d' < d$
 - Split the bucket, no need of directory doubling

Extendible Hashing

- Advantages:
 - Performance of file does not degrade with increase in size
 - Space overhead for directory is negligible
 - Splitting causes minor reorganization
- Disadvantages:
 - Directory must be searched before bucket. 2 block access instead of 1

Linear hashing

- Motivation: extendible hashing needs directory which doubles
- Q: can we do something simpler, with smoother growth?
- A: split buckets from left to right, regardless of which one overflowed

Initially: $h(x) = x \bmod N$ ($N=3$ here)

Assume capacity: 2 records / bucket


Use two hash functions:

$h_0(x) = x \bmod N$ ($N=3$ here) - for unsplit buckets

$h_1(x) = x \bmod (2*N)$ ($N=3$ here) - for the splitted ones


Linear Hashing Example

- In the following $M=3$ (initial # of buckets)
- Each bucket has 2 keys. One extra key for overflow.
- s is a pointer, pointing to the split location. This is the place where next split should take place.
- Insert Order: 1,7,3,8,12,4,11,2,10,13
- After insertion till 12:




Linear Hashing example

- When 4 inserted overflow occurred. So we split the bucket (no matter it is full or partially empty). And increment pointer.




- split bucket 0 and rehashed all keys.
- Placed 3 to new bucket as $(3 \bmod 6 = 3)$ and $(12 \bmod 6 = 0)$.
- Then 11 and 2 are inserted.
- s is pointing to bucket 1, hence split bucket 1 by re- hashing it.




Linear Hashing Example

- After split:




- Insertion of 10:
 - $(10 \bmod 3 = 1)$ and bucket $1 < s$, we need to hash 10 again using $h_1(10) = 10 \bmod 6 = 4^{\text{th}}$ bucket
- For 13
 - same bucket
 - Overflow
 - split 2nd bucket.



Linear Hashing Example

- Final Hash Table:



- s is moved to the top again as one cycle is completed and *level* is incremented.

Linear hashing - Searching

- Algo to find key 'k':

 compute $b = h_0(k)$ // original slot

 if $b < s$ // has already split

 compute $b = h_1(k)$

 search bucket b

Linear hashing - Deletion

- Inverse of insertion
- If underflow, **contract**
- If the last bucket is empty,
 - remove it and
 - Decrement s .
- If s is 0 and the last bucket becomes empty,
 - s is made to point to bucket $(n/2)-1$, where n is the current number of buckets,
 - Level is decremented, and
 - the empty bucket is removed.

B+ Trees vs Hashing

B+ Trees	Hashing
<ul style="list-style-type: none">• Speed on Search<ul style="list-style-type: none">• Exact match queries, worst case• Range queries• Nearest-neighbor queries• Speed on insertion + deletion• Smooth growing and shrinking (no-reorg)	<ul style="list-style-type: none">• Speed<ul style="list-style-type: none">• On exact match queries, on the average