

# MC 302 – DBMS: Concurrency Control 2

Goonjan Jain

Dept. of Applied Mathematics

Delhi Technological University

# Timestamp Allocation

- Each transaction  $T_i$  is assigned a unique fixed timestamp that is monotonically increasing.
  - Let **TS**( $T_i$ ) be the timestamp allocated to transaction  $T_i$
  - Different schemes assign timestamps at different times during the transaction.
- Multiple implementation strategies:
  - System Clock.
  - Logical Counter.
  - Hybrid.

# Timestamp Ordering Concurrency Control

- Use these timestamps to determine the serializability order.
- If **TS(Ti) < TS(Tj)**, then the DBMS must ensure that the execution schedule is equivalent to a serial schedule where Ti appears before Tj.

# Basic T/O

- Transactions read and write objects without locks.
- Every object  $X$  is tagged with timestamp of the last transaction that successfully did read/write:
  - **W-TS( $X$ )** – Write timestamp on  $X$
  - **R-TS( $X$ )** – Read timestamp on  $X$
- Check timestamps for every operation:
  - If transaction tries to access an object “from the future”, it aborts and restarts.

# Basic T/O – Reads and Writes

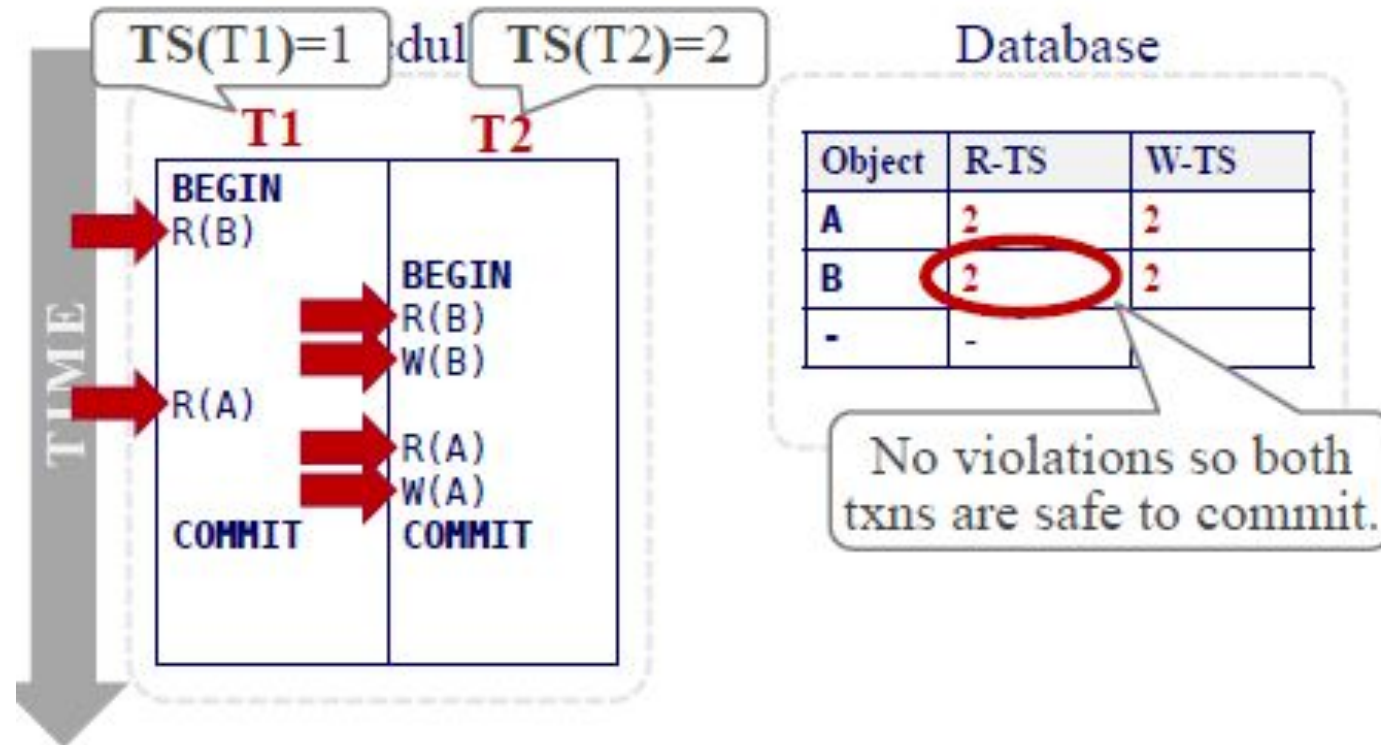
- Reads

- If  $\mathbf{TS(Ti)} < \mathbf{W-TS(X)}$ , this violates timestamp order of Ti w.r.t. writer of X.
  - Abort Ti and restart it (with same TS? why?)
- Else:
  - Allow Ti to read X.
  - Update  $\mathbf{R-TS(X)}$  to  $\mathbf{\max(R-TS(X), TS(Ti))}$
  - Have to make a local copy of X to ensure repeatable reads for Ti.

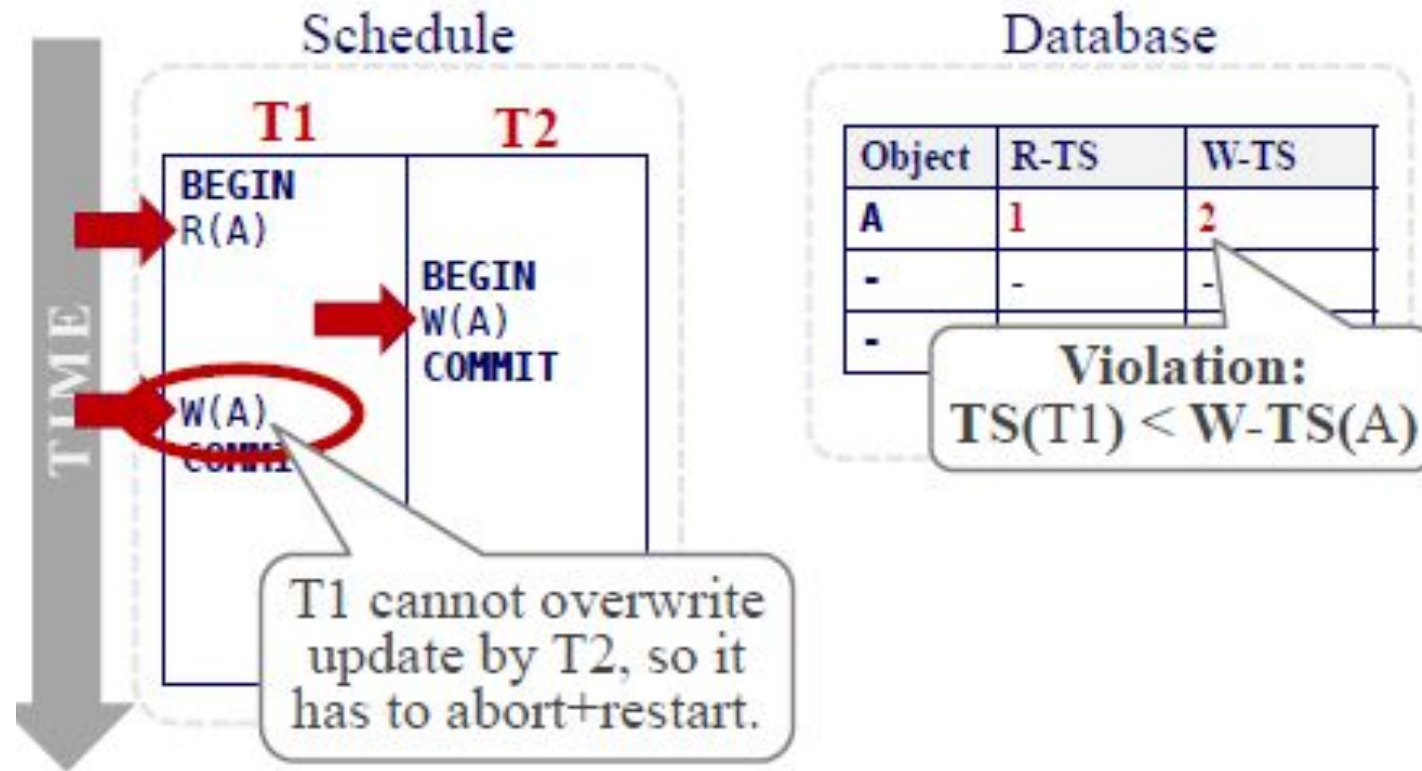
- Writes

- If  $\mathbf{TS(Ti)} < \mathbf{R-TS(X)}$  or  $\mathbf{TS(Ti)} < \mathbf{W-TS(X)}$ 
  - Abort and restart Ti.
- Else:
  - Allow Ti to write X and update  $\mathbf{W-TS(X)}$
  - Also have to make a local copy of X to ensure repeatable reads for Ti.

# Basic T/O Example



# Basic T/O Example 2

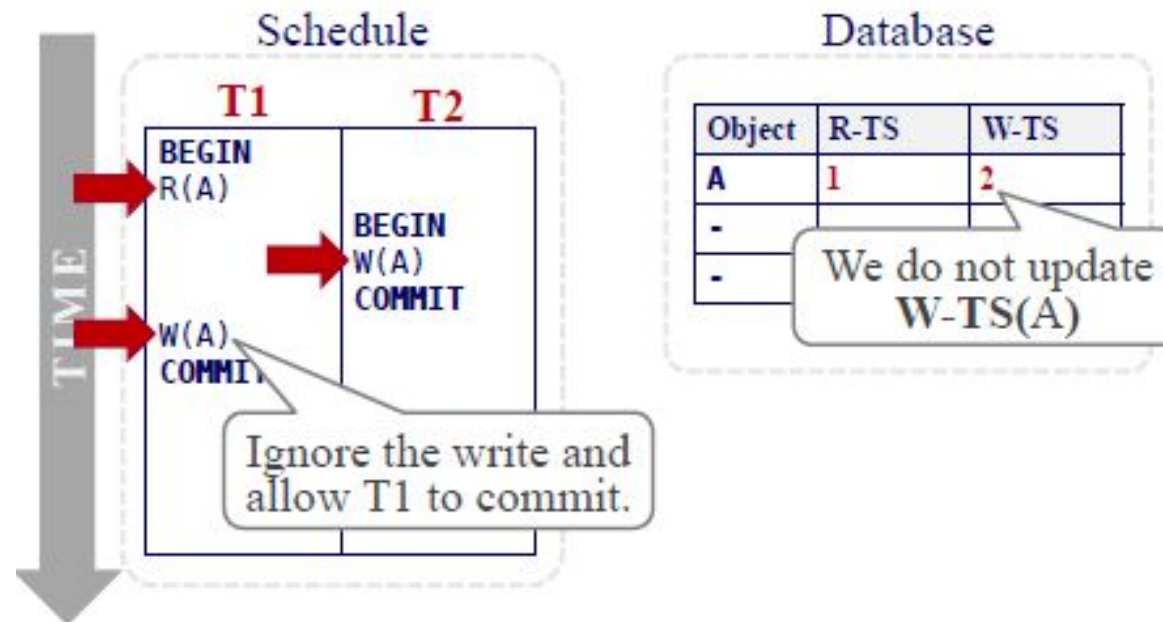


# Basic T/O: Thomas Write Rule

- If  **$TS(T_i) < R-TS(X)$** :
  - Abort and restart  $T_i$ .
- If  **$TS(T_i) < W-TS(X)$** :
  - **Thomas Write Rule**: Ignore the write and allow the txn to continue.
  - This violates timestamp order of  $T_i$
- Else:
  - Allow  $T_i$  to write  $X$  and update  **$W-TS(X)$**



# Basic T/O: Thomas Write Rule

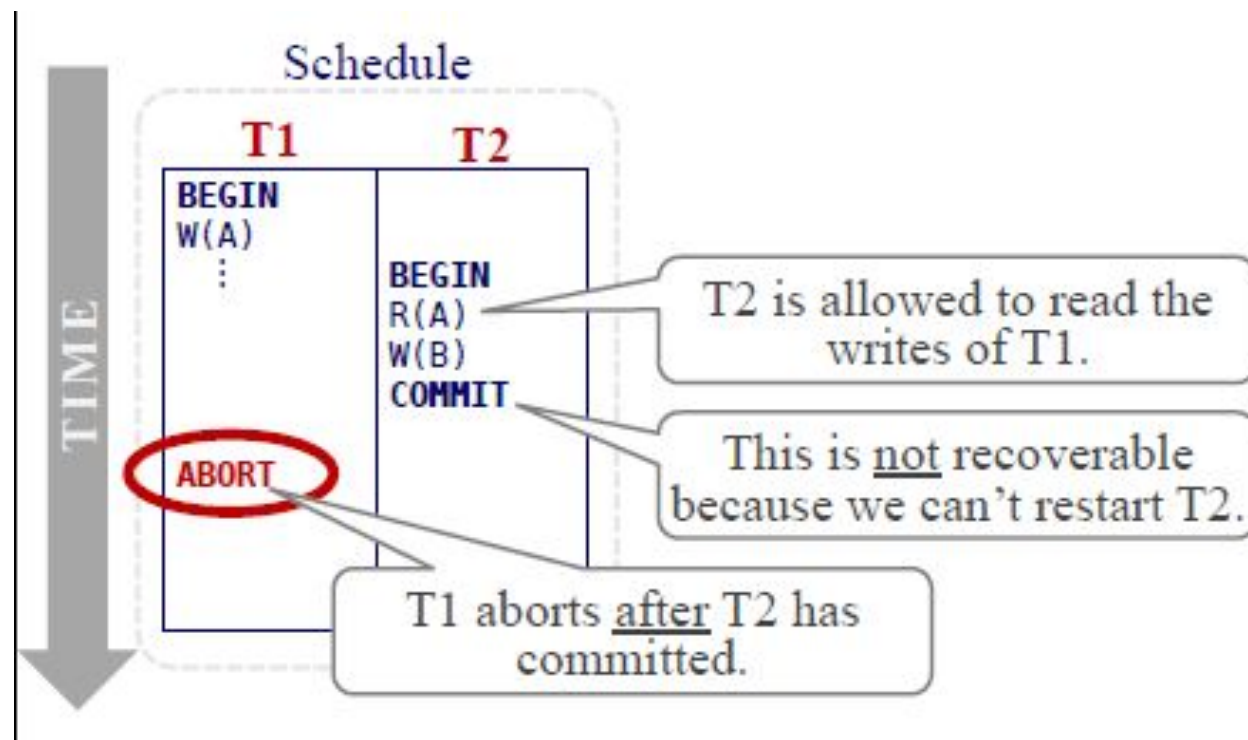


# Basic T/O

- Ensures conflict serializability if you don't use the Thomas Write Rule.
- No deadlocks because no transaction ever waits.
- Possibility of starvation for long transactions if short transactions keep causing conflicts.
- Permits schedules that are not ***recoverable***.

# Recoverable Schedules

- Transactions commit only after all transactions whose changes they read, commit.



# Basic T/O: Performance Issues

- High overhead from copying data to transaction's workspace and from updating timestamps.
- Long running transactions can get starved.
- Suffers from timestamp bottleneck.

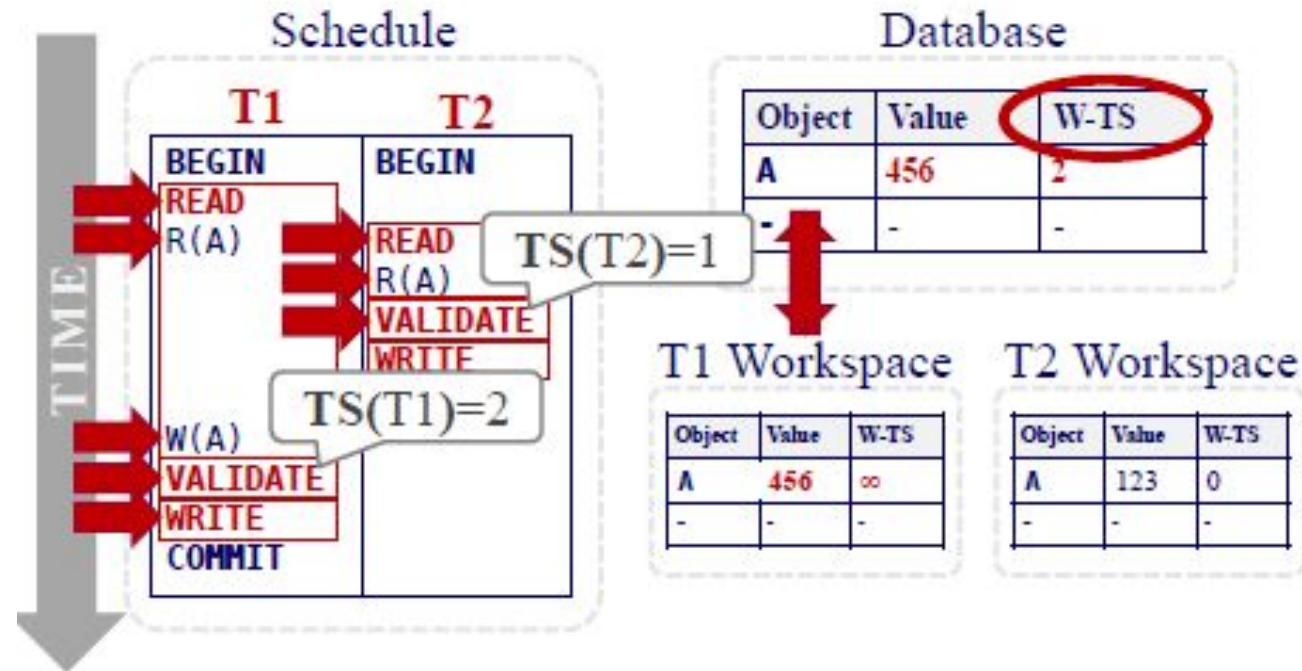
# Optimistic Concurrency Control

- Assumption: Conflicts are rare
- Forcing transactions to wait to acquire locks adds a lot of overhead.
- Optimize for the no-conflict case.

## OCC Phases:

- **Read**: Track the read/write sets of transactions and store their writes in a private workspace.
- **Validation**: When a transaction commits, check whether it conflicts with other transactions.
- **Write**: If validation succeeds, apply private changes to database. Otherwise abort and restart the transaction.

# OCC Example



# OCC: Validation Phase

- Need to guarantee only serializable schedules are permitted.
- At validation,  $T_i$  checks other txns for RW and WW conflicts and makes sure that all conflicts go one way (from older txns to younger txns).

# OCC: Serial Validation

- Maintain global view of all active transactions.
- Record read set and write set while transactions are running and write into private workspace.
- Execute **Validation** and **Write** phase inside a protected critical section.




# OCC: Validation Phase

- Each transaction's timestamp is assigned at the beginning of the validation phase.
- Check the timestamp ordering of the committing transaction with all other running transactions.
- If  $TS(T_i) < TS(T_j)$ , then one of the following three conditions must hold...

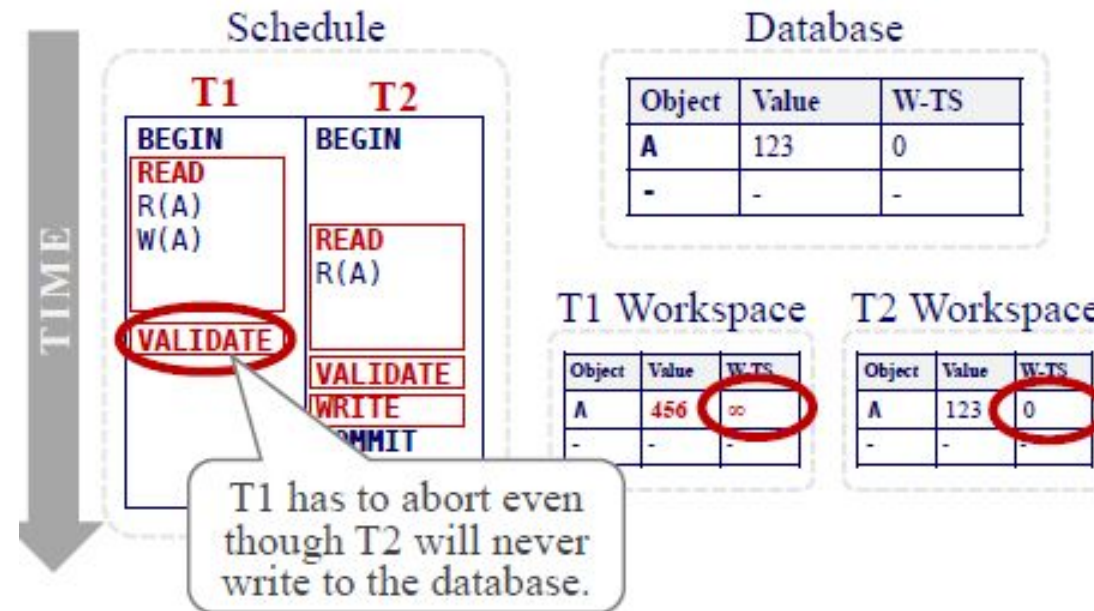
# OCC Validation #1

- $T_i$  completes all three phases before  $T_j$  begins.

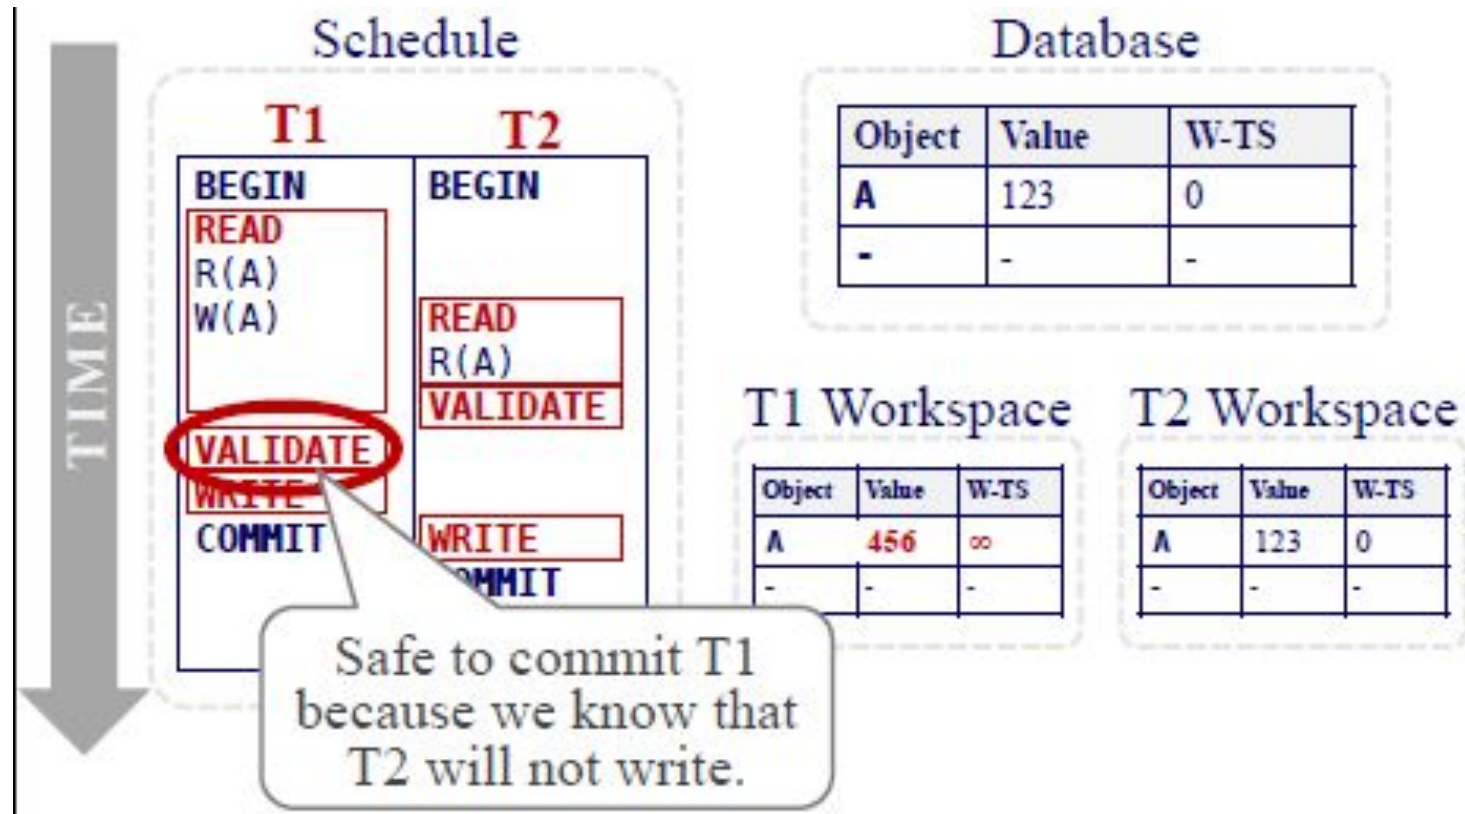


# OCC – Validation#2

- $T_i$  completes before  $T_j$  starts its **Write** phase, and  $T_i$  does not write to any object read by  $T_j$ .
  - $\text{WriteSet}(T_i) \cap \text{ReadSet}(T_j) = \emptyset$

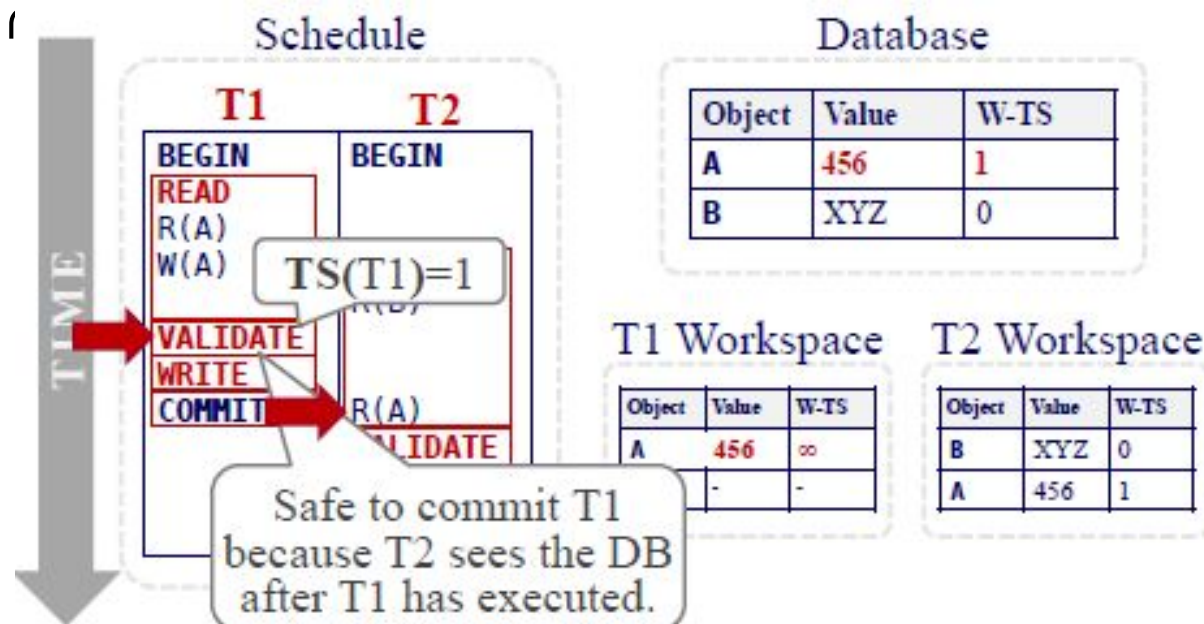


# OCC-Validation #2



# OCC-Validation#3

- $T_i$  completes its **Read** phase before  $T_j$  completes its **Read** phase
- And  $T_i$  does not write to any object that is either read or written by  $T_j$ :
  - $WriteSet(T_i) \cap ReadSet(T_j) = \emptyset$
  - $WriteSet(T_i) \cap WriteSet(T_j) = \emptyset$



# OCC-Observations

- **Q:** When does OCC work well?
- **A:** When # of conflicts is low:
  - All transactions are read-only (ideal).
  - Transactions access disjoint subsets of data.
- If the database is large and the workload is not skewed, then there is a low probability of conflict, so again locking is wasteful.

# OCC Performance Issues

- High overhead for copying data locally.
- **Validation/Write** phase bottlenecks.
- Aborts are more wasteful because they only occur *after* a txn has already executed.
- Suffers from timestamp allocation bottleneck.

# Multi-Version Concurrency Control

- Writes create new versions of objects instead of in-place updates:
  - Each successful write results in the creation of a new version of the data item written.
- Use write timestamps to label versions.
  - Let  $X_k$  denote the version of  $X$  where for a given transaction  $T_i$ :  $\mathbf{W-TS}(X_k) \leq \mathbf{TS}(T_i)$



# MVCC – Reads and Writes

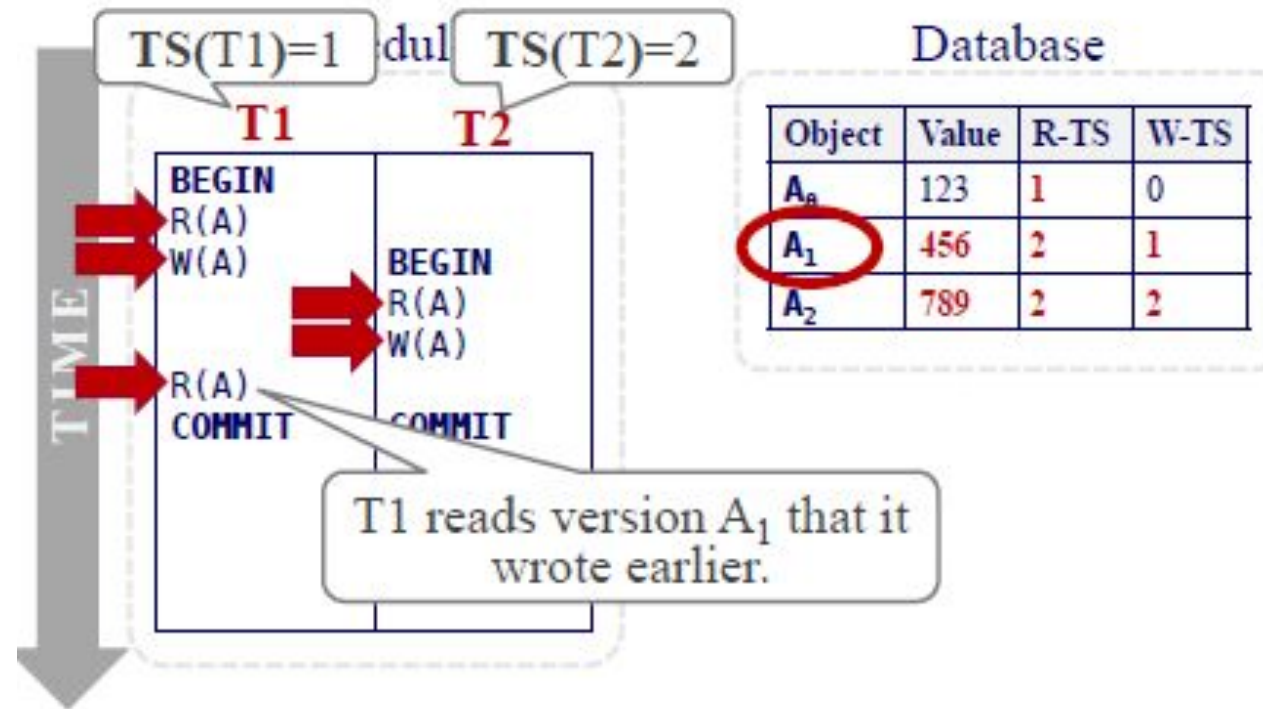
## Reads

- Any read operation sees the latest version of an object from right before that transaction started.
- Every read request can be satisfied without blocking the transaction.
- If  $TS(T_i) > R-TS(X_k)$ :
  - Set  $R-TS(X_k) = TS(T_i)$

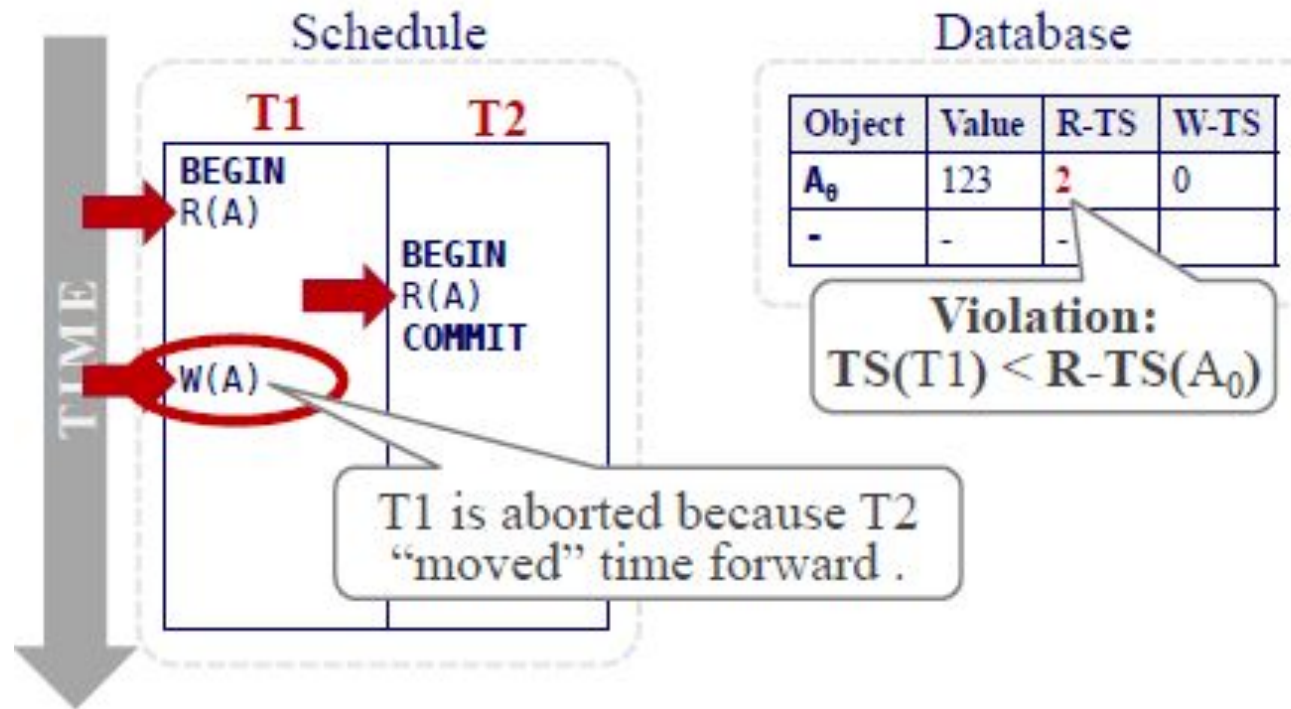
## Writes

- If  $TS(T_i) < R-TS(X_k)$ :
  - Abort and restart  $T_i$ .
- If  $TS(T_i) = W-TS(X_k)$ :
  - Overwrite the contents of  $X_k$ .
- Else:
  - Create a new version of  $X_{k+1}$  and set its write timestamp to  $TS(T_i)$ .

# MVCC – Example #1



# MVCC – Example #2



# MVCC

- Can still incur cascading aborts because a transaction sees uncommitted versions from transactions that started before it did.
- Old versions of tuples accumulate.
- The DBMS needs a way to remove old versions to reclaim storage space.