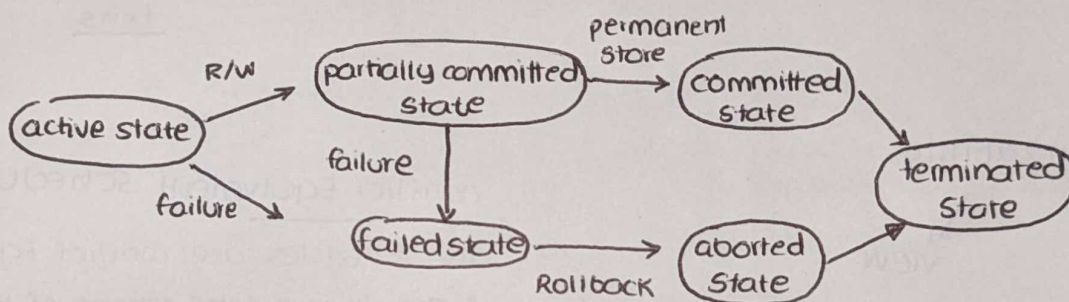
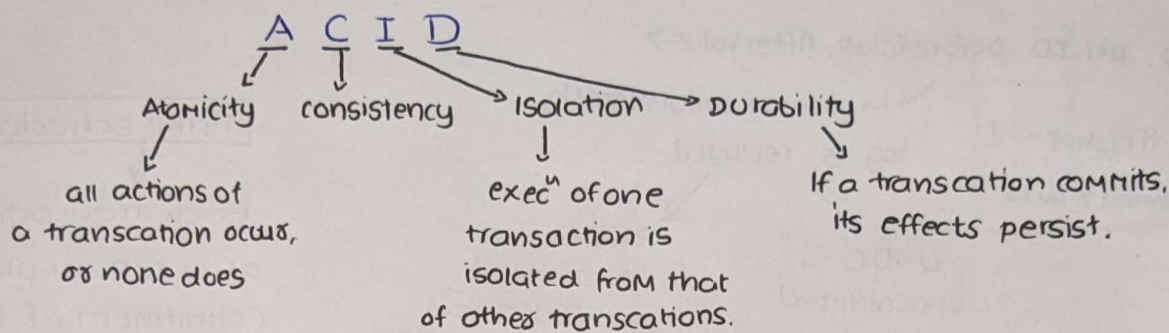


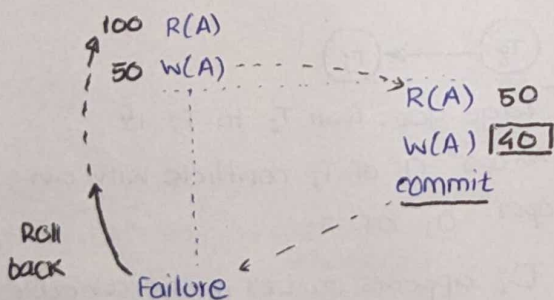
## # Transaction correctness Criteria :



## # Concurrency Problems

Dirty Read / Read After Write

T1                      T2



post failure, A becomes 100, but T2 has already worked on a wrong value of A = 50.

last update

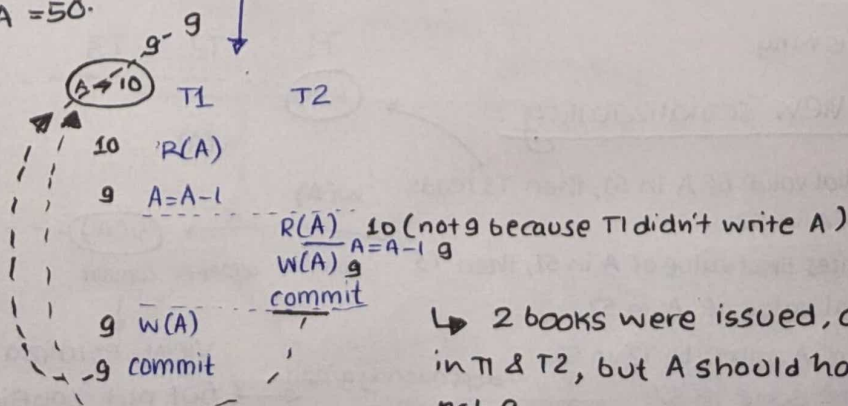
One commit overwrites the changes made by some prev. commits.

T1      T2  
- R(A) - R(A) → no problem

R(A)    W(A) → Read-Write Problem / Unrepeatable Read.

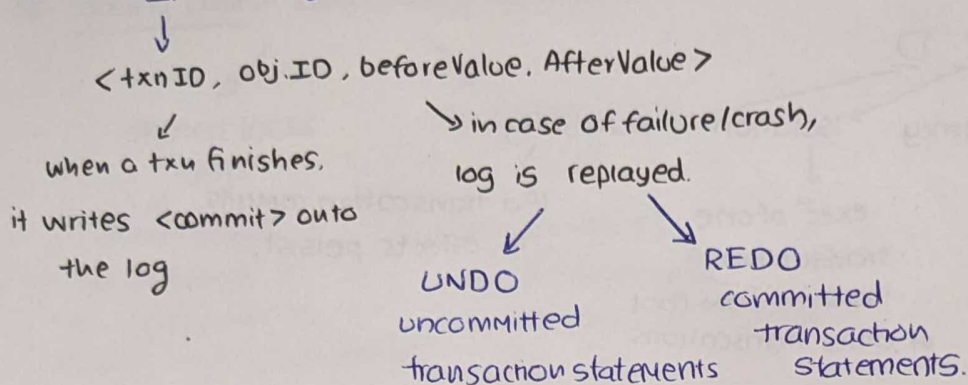
W(A)    R(A) → Dirty Read Problem / Read After Write.

W(A)    W(A)



2 books were issued, one each in T1 & T2, but A should have been 8, not 9.


## # Write-ahead log.



## serial schedules

Each transaction occurs  
after the completion/  
commitment of the prev.  
txn. No interleaving of  
txns.

## # Serializability



schedule S is conflict serializable  
iff

- S is conflict equivalent to some serial schedule
- Able to transform S into a serial sched. by swapping consecutive non-conflicting operations of different transactions.


## conflict Equivalent Schedules

Two schedules are conflict Equivalent if:

- They involve same actions of the same txns.
- Every pair of conflicting actions is ordered the same way.

another way to check conflict serializability

Make a directed graph (dependency graph)



Edge goes from  $T_i$  to  $T_j$  if

- An oper<sup>n</sup>  $O_i$  of  $T_i$  conflicts with an oper<sup>n</sup>  $O_j$  of  $T_j$
- $O_i$  appears earlier in the schedule than  $O_j$  does.

schedule is  
conflict serializable (and  
hence serializable) if  
there is no loop in the graph.

maybe  
serializable

if a loop  
exists

may not be  
serializable.


this is why

## view serializability

• If  $T_1$  reads initial value of A in  $S_1$ , then  $T_1$  reads initial value of A in  $S_2$ .

• If  $T_1$  reads writes final value of A in  $S_1$ , then  $T_1$  writes the final value of A in  $S_2$ .

★ •  $T_1$  reads value of A written by  $T_2$  in  $S_1$ , then  $T_1$  does the same in  $S_2$ .



view serializable

dependency graph  
is cyclic. (but not conflict serializable)



# # locking based concurrency control

shared locks  
S(A)

- (Read only)
- can be shared, i.e. can be accessed by multiple txns simultaneously.

Exclusive locks

- X(A)
- (Read + Write)
- can't be shared until one txn releases the lock.

compatibility matrix

|        |   | T2 wants |    |
|--------|---|----------|----|
| T1 has | S | S        | X  |
|        | X | yes      | no |
|        |   | X        | no |

Example: Last update w/o locks

| T1     | T2     |
|--------|--------|
| R(A)   |        |
| R(A)   |        |
| W(A)   |        |
| commit |        |
|        | W(A)   |
|        | commit |

conflict

with locks

| T1        | T2   |
|-----------|--|
| lock(A)   |  |
| R(A)      |  |
|           | lock(A) → denied because T1 has A's lock       |
| W(A)      |  |
| unlock(A) |  |
| commit    |  |
|           | R(A) → lock granted to A after T1 releases it. |
|           | W(A)   |
|           | unlock(A)                                      |
|           | commit   |

drawbacks

- may not always give serializable schedules
- may not be free from irrecoverability
- starvation
- deadlock

2PL (Phase locking)

growing

locks are acquired, no lock is released

shrinking

no lock is acquired, locks are released

follow compatibility matrix

If T1 is in growing phase, then T2 can acquire a shared lock

always ensures serializability

| T1               | T2              |
|------------------|-----------------|
| X(A)             |                 |
| R(A)             |                 |
| W(A)             |                 |
| U(A)             | S(A)            |
|                  | R(A)            |
|                  | commit          |
| rollback failure | can't roll back |

read after write (dirty read)

drawbacks

- May not be recoverable
- deadlock
- starvation
- cascading rollback/abort

sol<sup>n</sup>

strict 2PL

Rigorous 2PL

basic 2PL + all exclusive locks should hold until commit/abort.


basic 2PL + all X, S locks should hold until commit/abort.

## # Timestamp Ordering Protocol

$TS(T_i)$

timestamp assigned to  $i^{th}$  transaction.

not necessarily the time at which the txn enters; a random integer value.



Read Timestamp (RTS)

TS of the latest txn that performed a successful read operation.

Write Timestamp (WTS)

TS of the latest txn that performs a successful write op<sup>n</sup>.

### Basic T/O Rules

1) Transaction  $T_i$  issues a Read(A) op<sup>n</sup>:

a) if  $WTS(A) > TS(T_i) \Rightarrow$  Rollback  $T_i$

b) otherwise, execute  $R(A)$ ;

set  $RTS(A) = \max\{RTS(A), TS(T_i)\}$

2) Transaction  $T_i$  issues write(A) op<sup>n</sup>:

a) if  $RTS(A) > TS(T_i) \Rightarrow$  Rollback  $T_i$

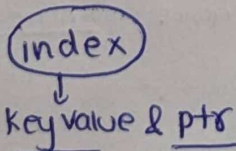
b) if  $WTS(A) > TS(T_i) \Rightarrow$  Rollback  $T_i$

c) otherwise, execute  $W(A)$  op<sup>n</sup>;

set  $WTS(A) = TS(T_i)$

Dense Index

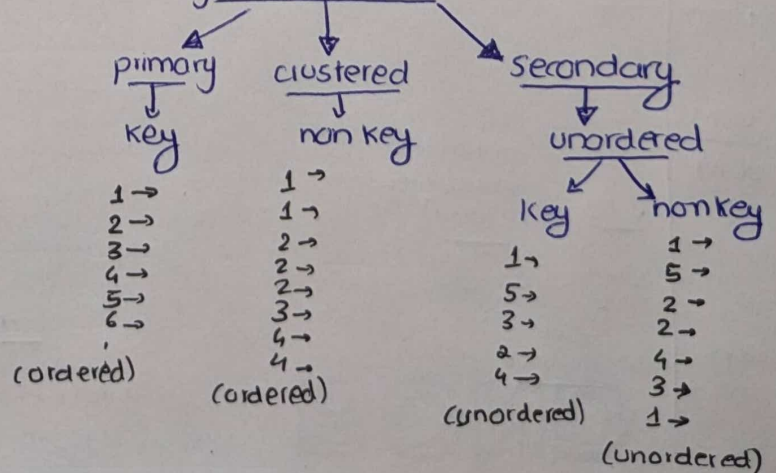
an entry for each key value of every record.



Sparse Index

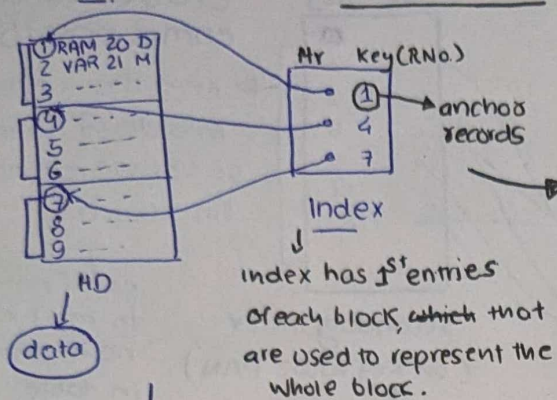
entries for only some of the key values out of all records

single level ordered indices





## # primary\_index → (sparse Index)



disadvantages: To insert a record at its correct pos<sup>n</sup>, move records to make space, we might change the anchor records in the data which may lead to an obsolete index, if not changed simultaneously, same for del<sup>n</sup>.

$$\text{No. of entries in Index} = \text{No. of blocks in memory}$$

search time:  $\log_2(N+1)$  searching for specific data after reaching the req. block.

Examples:

$\gamma = 30,000$  records

BlockSize(B) = 1024 B

RecordSize(R) = 100 B

Blocking factor =  $\frac{B}{R} = \frac{1024}{100} = 10$  records/block (BFR)

No. of blocks (b) =  $\frac{\gamma}{\text{BFR}} = \frac{30000}{10} = 3000$

using indexing:

field value size: 9 B  
Block ptr size: 6 B } 15 B (size of each entry)

Blocking factor for index =  $\frac{1024}{15} = 68$  entries/block

Total no. of index entries = 3000 (=no. of blocks)

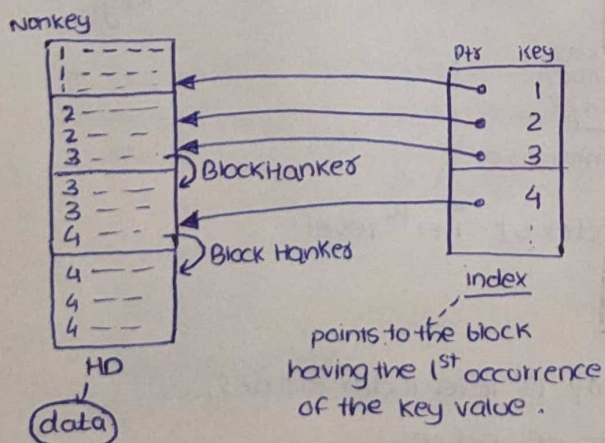
no. of index blocks  $b_i = \frac{3000}{68} = 45$

binary search on index requires  $\log_2 45 = 6$  block access

& searching a record using index =  $6 + 1 = 7$  block access

search time  $\approx \log_2 3000 + 1 = 12 + 1$  no. of block access. (without indexing)

## # clustering Index → (sparse Index)



rest is same as primary index

access time using index:

$\log_2 b_i + 1 + 1$   
bin. search on index    searching in the block    the entry might be present in the subsequent blocks.

disadvantages: insertion & deletion.

possible sol<sup>n</sup>:

reserve the whole block for each value of a clustering index.



## # Secondary Index (on key value)

Unordered (on Non Key Value)

| EID | Name | PAN |
|-----|------|-----|
| 1   | A    | 40  |
| 2   | B    | 51  |
| 3   | A    | 63  |
| 4   | C    | 21  |
| 5   | B    | 5   |
| 6   | E    | 18  |
| 7   | F    | 1   |
| 8   | A    | 6   |
| 9   | G    | 2   |

sparse index ← primary index (based on EID)

fails if the results are to be searched on the basis of "PAN"

Primary key (Ordered)  
non key (Unordered)  
Candidate key (Unordered)

dense index (no blocks anchor entries possible)

keep the keys in ordered manner so that we can apply bin. search.

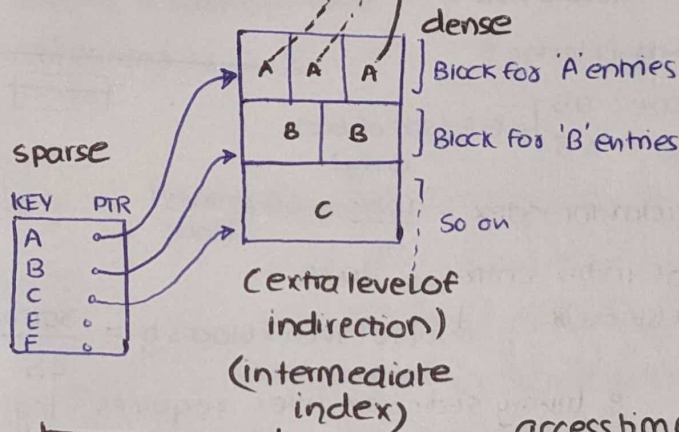
secondary Index (on key value 'PAN')

access time =

$$\log_2 \frac{N}{L} + 1$$

no. of entries in index = no. of entries in table.

no. of entries (total)



secondary Index (on Non key value 'Name')


access time

$$= \log_2 N + \frac{N}{L} + \frac{1}{L} \text{ retrieve data}$$

# 1<sup>st</sup> index entries

No. of entries in block of a particular key

## # Multilevel Index



$$\text{Fanout}(F_0) = \frac{B}{P+V}$$

Block size

pointer size

key value size

1<sup>st</sup> level index entry size

• Number of entries in the index at  $i+1$ <sup>th</sup> level

$$\delta_{i+1} = \lceil \delta_i / F_0 \rceil$$

• A multilevel index with  $\delta_i$  1<sup>st</sup> level index entries, will have approx. 't' levels of indices,

$$t = \lceil \log_{F_0} \delta_i \rceil$$

## # Multilevel Dynamic Indexing Using B-Trees (Balanced Trees)

- $p \rightarrow$  Max no. of children a node can have
- Each node, except leaf and roots, has atleast  $\lceil p/2 \rceil$  pointers (children)
- All leaf nodes are at the same level. All the pointers of the leaf nodes are NULL.

